



INTRODUCTION TO DATA ANALYTICS

Linear Regression

Asst. Prof. Dr. Rathachai Chawuthai

Department of Computer Engineering

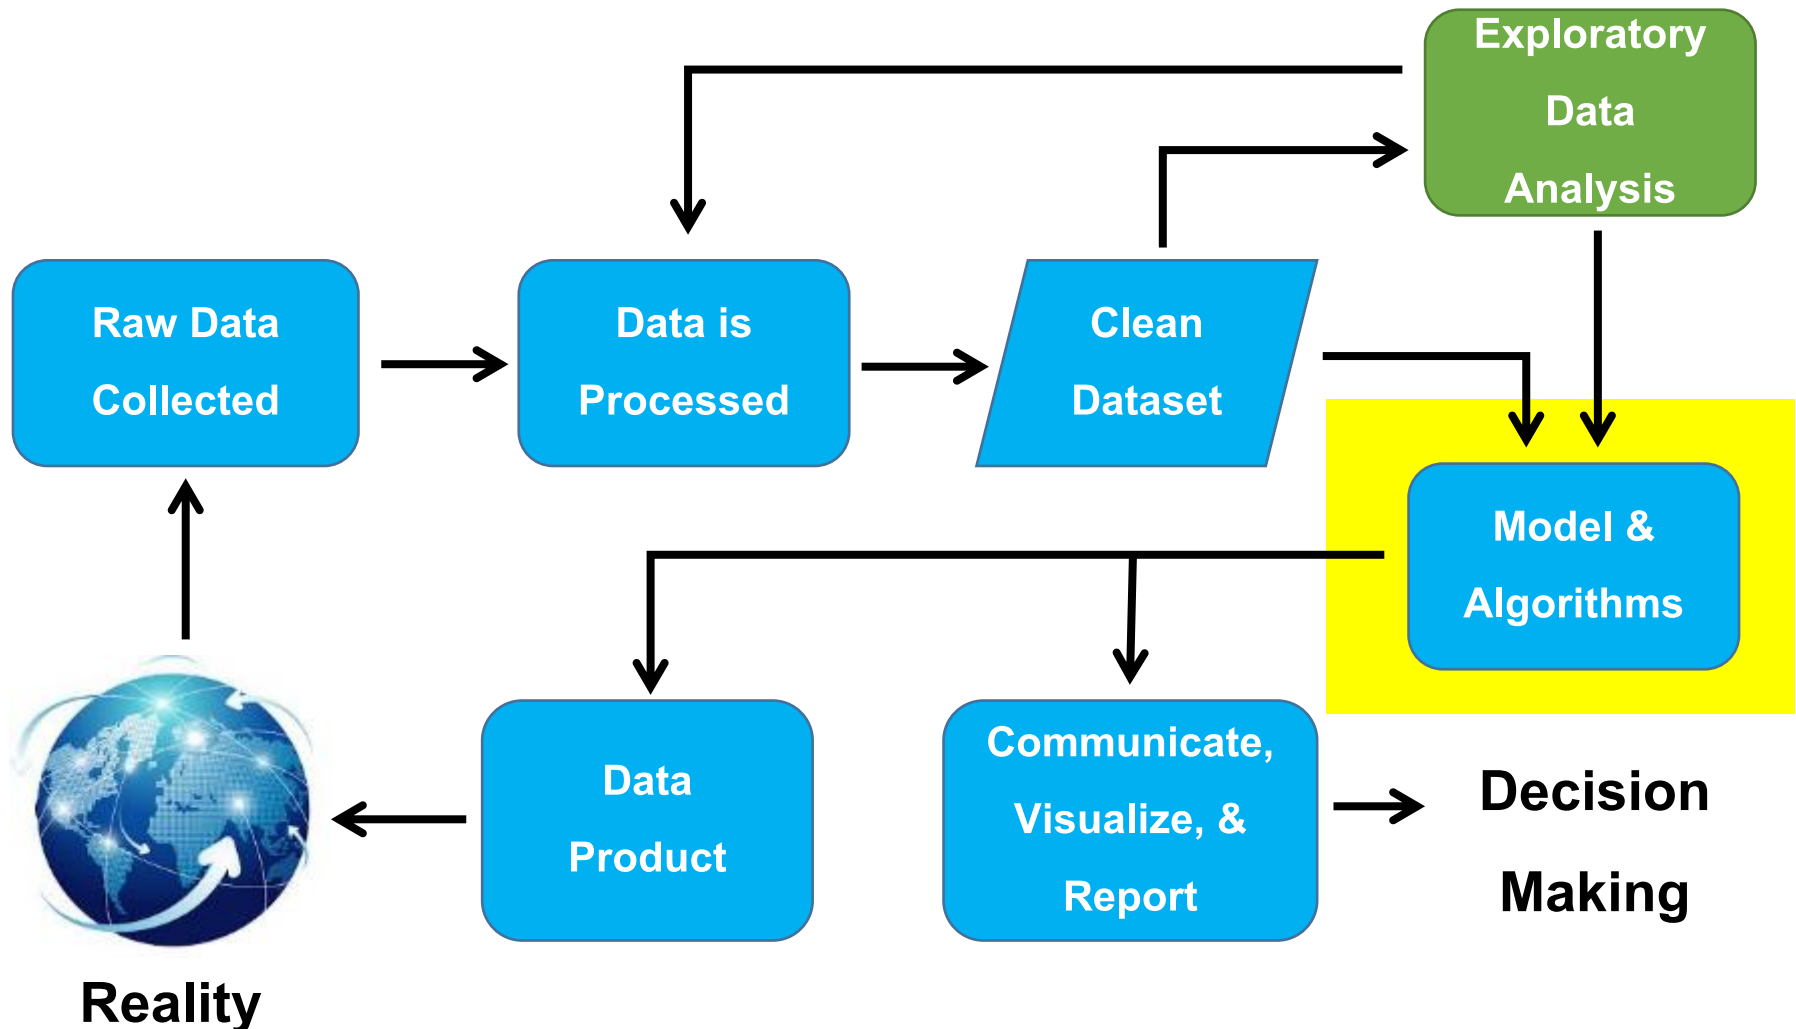
Faculty of Engineering

King Mongkut's Institute of Technology Ladkrabang

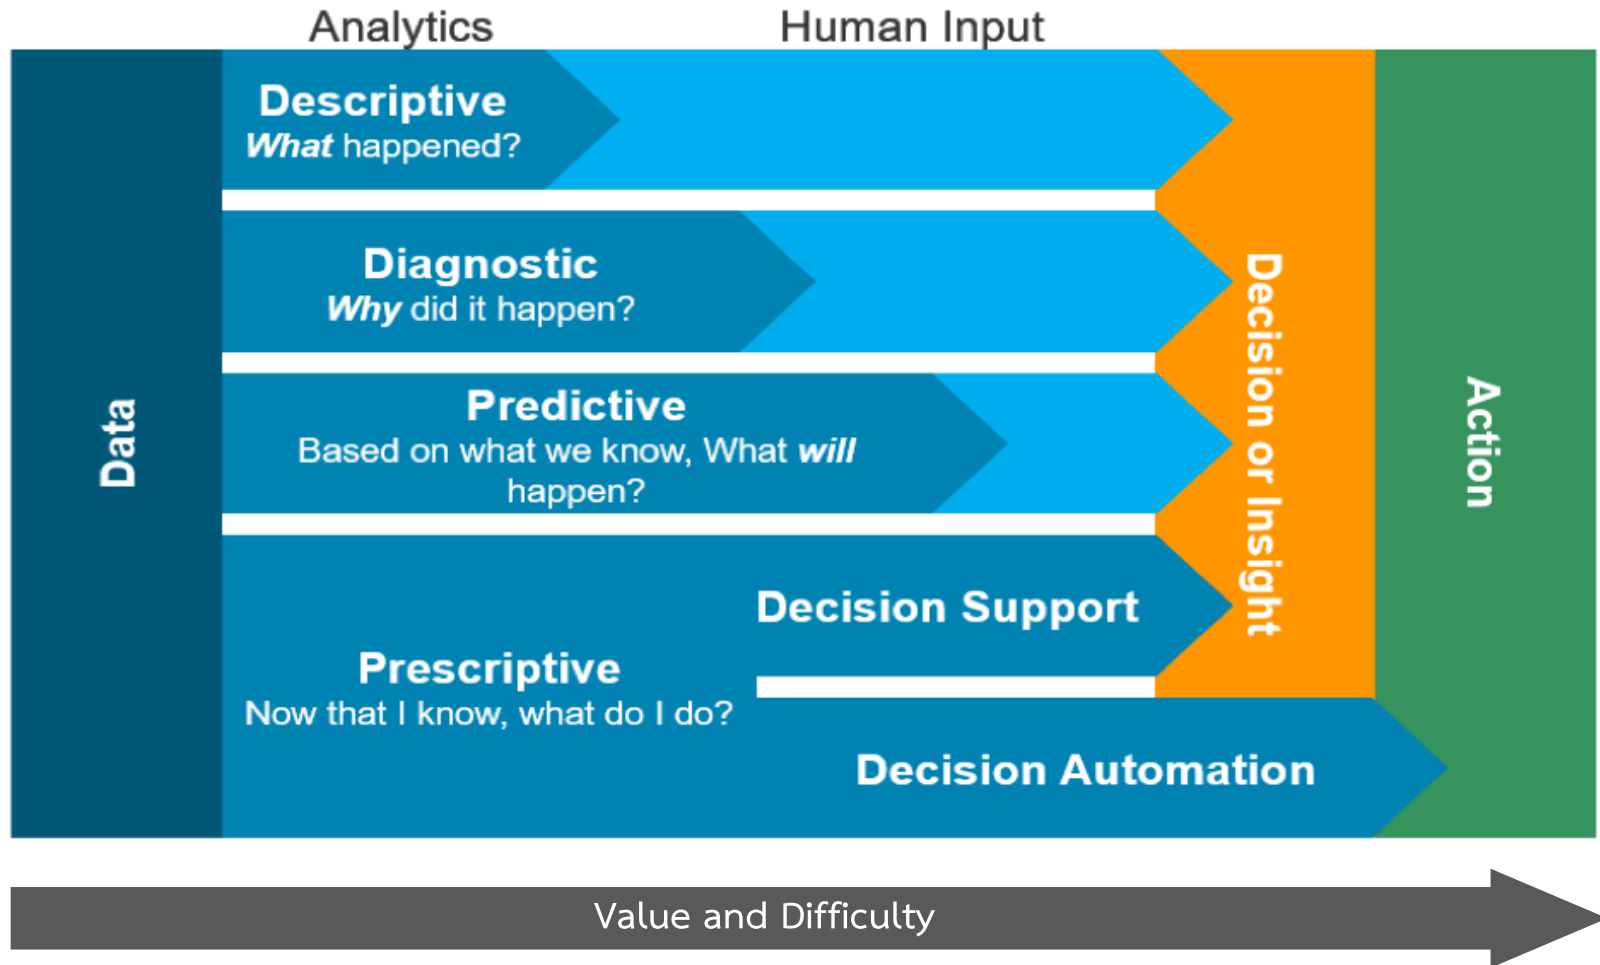
Agenda

- Linear Regression
- Evaluation Methods
- Feature Selection
- Cross Validation
- Feature Scaling

Data Science Process



Data Analytics



Machine Learning

```
graph TD; ML[Machine Learning] --> SL[Supervised Learning]; ML --> UL[Unsupervised Learning]; SL --> R[Regression]; SL --> C[Classification]; UL --> Cl[Clustering]; R --- RList["• Linear Regression<br>• Polynomial Regression"]; C --- CList["• Decision Tree<br>• Logistic Regression<br>• Neural Network<br>• etc."]; Cl --- ClList["• K-Means<br>• DB-SCAN<br>• etc."];
```

Supervised Learning

Develop predictive model based on both input and output data

Regression

- Linear Regression
- Polynomial Regression

Classification

- Decision Tree
- Logistic Regression
- Neural Network
- etc.

Unsupervised Learning

Develop predictive model based on both input and output data

Clustering

- K-Means
- DB-SCAN
- etc.

Linear Regression

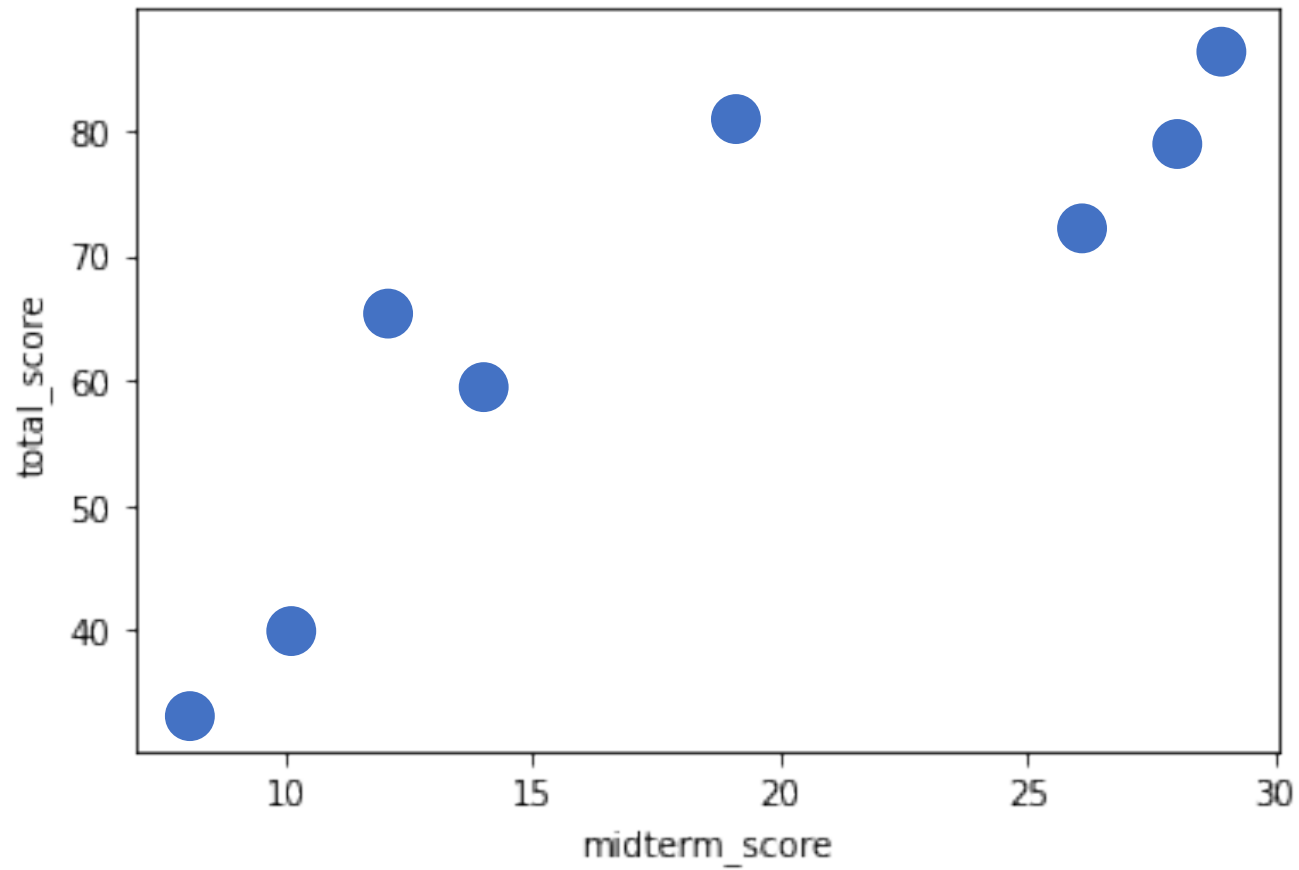


Course Score


midterm_score	total_score
10	40
8	33
14	59
12	65
19	81
26	72
28	79
29	87

(assumed data)


Prediction?




Prediction?




Prediction?



Prediction?



Linear Equation



$$y = ax + c$$

Red arrows point from the question marks in the equation to the corresponding parts of the graph: the first arrow points from the coefficient 'a' to the slope of line A, and the second arrow points from the constant 'c' to the y-intercept of line C.

Simple Linear Regression

$$\hat{Y}_i = b_0 + b_1 X_i$$

Simple Linear Regression

$$\hat{Y}_i = b_0 + b_1 X_i$$

(actual values)

total_score
40
33
59
65
81
72
79
87

total_score

midterm_score
10
8
14
12
19
26
28
29

Course Score

midterm_score	hour_study	total_score
10	3	40
8	2	33
14	5	59
12	5	65
19	7	81
26	8	72
28	8	79
29	9	87

(assumed data)

Multiple Linear Regression

$$\hat{Y}_i = b_0 + b_1 X_{1i} + b_2 X_{2i} + \dots$$

Multiple Linear Regression

$$\hat{Y}_i = b_0 + b_1X_{1i} + b_2X_{2i} + \dots$$

actual	total_score
40	
33	
59	
65	
81	
72	
79	
87	

midterm_score
10
8
14
12
19
26
28
29

hour_study
3
2
5
5
7
8
8
9

Coefficients

$$\hat{Y}_i = 19.8 - 1.8 * \textit{midterm_score} + 13.3 * \textit{hour_study}$$

Python

```
from sklearn import linear_model
X = df[["midterm_score", "hour_study"]]
y = df["total_score"]
lm = linear_model.LinearRegression()
lm.fit(X, y)
print("intercept=", lm.intercept_)
print("coef=", lm.coef_)
```




```
intercept= 19.796670104611792
coef= [-1.84647119  13.34492412]
```

Evaluation Methods



Residual



Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (\hat{y}_i - y_i)^2}{N}}$$

Diagram illustrating the RMSE formula with annotations:

- predicted**: Points to \hat{y}_i (predicted value).
- actual**: Points to y_i (actual value).
- number of samples**: Points to N (total number of samples).

RMSE

$$\hat{Y}_i = 19.8 - 1.8 * mid + 13.3 * hour$$

mid	hour	y_i
10	3	40
8	2	33
14	5	59
12	5	65
19	7	81
26	8	72
28	8	79
29	9	87

RMSE

$$\hat{Y}_i = 19.8 - 1.8 * mid + 13.3 * hour$$

mid	hour	y_i	\hat{y}_i
10	3	40	41.4
8	2	33	31.7
14	5	59	60.7
12	5	65	64.4
19	7	81	78.1
26	8	72	78.5
28	8	79	74.9
29	9	87	86.4

RMSE

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (\hat{y}_i - y_i)^2}{N}}$$

$$\hat{Y}_i = 19.8 - 1.8 * mid + 13.3 * hour$$

mid	hour	y_i	\hat{y}_i	$\hat{y}_i - y_i$	$(\hat{y}_i - y_i)^2$
10	3	40	41.4	1.4	1.9
8	2	33	31.7	-1.3	1.7
14	5	59	60.7	1.7	2.8
12	5	65	64.4	-0.6	0.4
19	7	81	78.1	-2.9	8.2
26	8	72	78.5	6.5	42.9
28	8	79	74.9	-4.1	17.2
29	9	87	86.4	-0.6	0.4

RMSE

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (\hat{y}_i - y_i)^2}{N}}$$

$$\hat{Y}_i = 19.8 - 1.8 * mid + 13.3 * hour$$

mid	hour	y_i	\hat{y}_i	$\hat{y}_i - y_i$	$(\hat{y}_i - y_i)^2$
10	3	40	41.4	1.4	1.9
8	2	33	31.7	-1.3	1.7
14	5	59	60.7	1.7	2.8
12	5	65	64.4	-0.6	0.4
19	7	81	78.1	-2.9	8.2
26	8	72	78.5	6.5	42.9
28	8	79	74.9	-4.1	17.2
29	9	87	86.4	-0.6	0.4
Sum					75.4
Mean					9.4
RMSE					3.1

Other Evaluation Methods

Mean squared error

$$\text{MSE} = \frac{1}{n} \sum_{t=1}^n e_t^2$$

Root mean squared error

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{t=1}^n e_t^2}$$

Mean absolute error

$$\text{MAE} = \frac{1}{n} \sum_{t=1}^n |e_t|$$

Mean absolute percentage error

$$\text{MAPE} = \frac{100\%}{n} \sum_{t=1}^n \left| \frac{e_t}{y_t} \right|$$

Python

1. `from sklearn import linear_model`
2. `from sklearn.metrics import mean_squared_error`
3. `from math import sqrt`
4. `X = df[“midterm_score”, “hour_study”]`
5. `y = df[“total_score”]`
6. `lm = linear_model.LinearRegression()`
7. `lm.fit (X,y)`
8. `y_predicted = lm.predict(X)`
9. `rmse = mean_squared_error(y, y_predicted, squared=False)`

Learning



Solve by Matrix

SOLVE

$$\begin{cases} 3x + 5y = -1 \\ x - 2y = 4 \end{cases} \quad A^{-1} = \frac{1}{ad-bc} \begin{pmatrix} d & -b \\ -c & a \end{pmatrix}$$
$$\begin{bmatrix} 3 & 5 \\ 1 & -2 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} -1 \\ 4 \end{bmatrix}$$
$$A^{-1} = \frac{1}{3(-2) - 1(5)} \begin{pmatrix} -2 & -5 \\ -1 & 3 \end{pmatrix}$$

Solve by



SOLVE $\begin{cases} 3x + 5y = -1 \\ x - 2y = 4 \end{cases} \quad A^{-1} = \frac{1}{ad-bc} \begin{pmatrix} d & -b \\ -c & a \end{pmatrix}$

$$\begin{bmatrix} 3 & 5 \\ 1 & -2 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} -1 \\ 4 \end{bmatrix}$$
$$A^{-1} = \frac{1}{3(-2) - 1(5)} \begin{pmatrix} -2 & -5 \\ -1 & 3 \end{pmatrix}$$



New algorithm ($n \geq 1$):

Repeat {


$$\downarrow \frac{\partial}{\partial \theta_j} J(\theta)$$

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$


(simultaneously update θ_j for $j = 0, \dots, n$) }

Learning based on Gradient Descent

$$y = \textcolor{red}{a}x + c$$




a




Learning


$$y = \textcolor{red}{a}x + c$$



a




Learning




a

$$y = \textcolor{red}{a}x + c$$




Learning

$$y = \textcolor{red}{a}x + c$$




a




Learning

$$y = \textcolor{red}{a}x + c$$




a




Learning

$$y = \textcolor{red}{a}x + c$$




a




Learning



$$y = \textcolor{red}{a}x + c$$




a



Learning



Learning




cost: $\sum |t - y|^2$





n-th round

Learning





Learning





n-th round

Learning





n -th round

Learning




n -th round



Learning




Learning



Learning






Gradient Descent for Linear Regression


$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

If α is too small, gradient descent can be slow.

If α is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge.



Gradient Descent for Linear Regression



Feature Selection




Boston Housing

CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV	CAT. MEDV
0.00632	18	2.31	0	0.538	6.575	65.2	4.09	1	296	15.3	396.9	4.98	24	0
0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.9	9.14	21.6	0
0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7	1
0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4	1
0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.9	5.33	36.2	1
0.02985	0	2.18	0	0.458	6.43	58.7	6.0622	3	222	18.7	394.12	5.21	28.7	0
0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.6	12.43	22.9	0
0.14455	12.5	7.87	0	0.524	6.172	96.1	5.9505	5	311	15.2	396.9	19.15	27.1	0
0.21124	12.5	7.87	0	0.524	5.631	100	6.0821	5	311	15.2	386.63	29.93	16.5	0
0.17004	12.5	7.87	0	0.524	6.004	85.9	6.5921	5	311	15.2	386.71	17.1	18.9	0
0.22489	12.5	7.87	0	0.524	6.377	94.3	6.3467	5	311	15.2	392.52	20.45	15	0
0.11747	12.5	7.87	0	0.524	6.009	82.9	6.2267	5	311	15.2	396.9	13.27	18.9	0
0.09378	12.5	7.87	0	0.524	5.889	39	5.4509	5	311	15.2	390.5	15.71	21.7	0
0.62976	0	8.14	0	0.538	5.949	61.8	4.7075	4	307	21	396.9	8.26	20.4	0
0.63796	0	8.14	0	0.538	6.096	84.5	4.4619	4	307	21	380.02	10.26	18.2	0

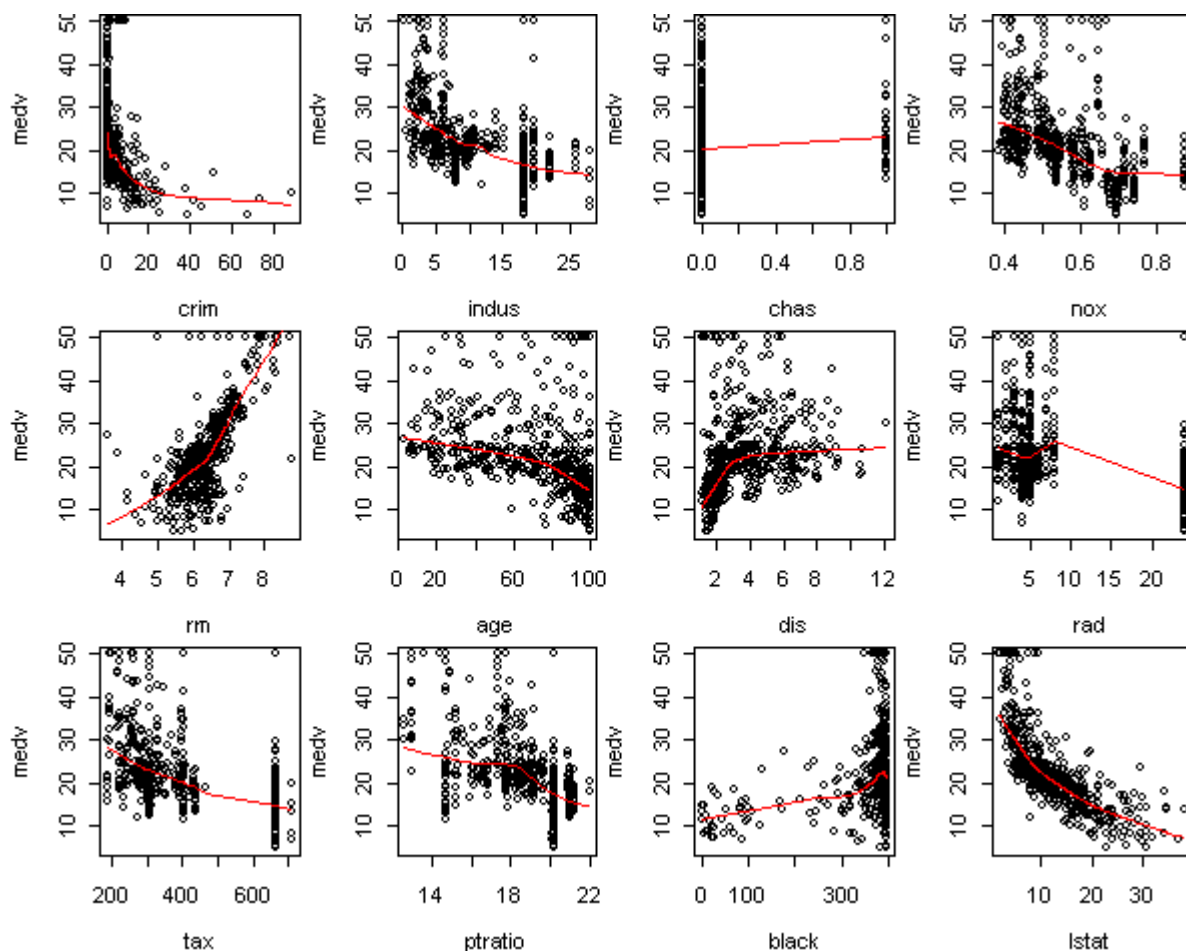
Boston Housing

CRIM	Per capita crime rate by town
ZN	Proportion of residential land zoned for lots over 25,000 sq.ft.
INDUS	Proportion of non-retail business acres per town
CHAS	Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
NOX	Nitric oxides concentration (parts per 10 million)
RM	Average number of rooms per dwelling
AGE	Proportion of owner-occupied units built prior to 1940
DIS	Weighted distances to five Boston employment centers
RAD	Index of accessibility to radial highways
TAX	Full-value property-tax rate per \$10,000
PTRATIO	Pupil-teacher ratio by town
B	$1000(B_k - 0.63)^2$ where B_k is the proportion of African-Americans by town
LSTAT	% Lower status of the population
MEDV	Median value of owner-occupied homes in \$1000's



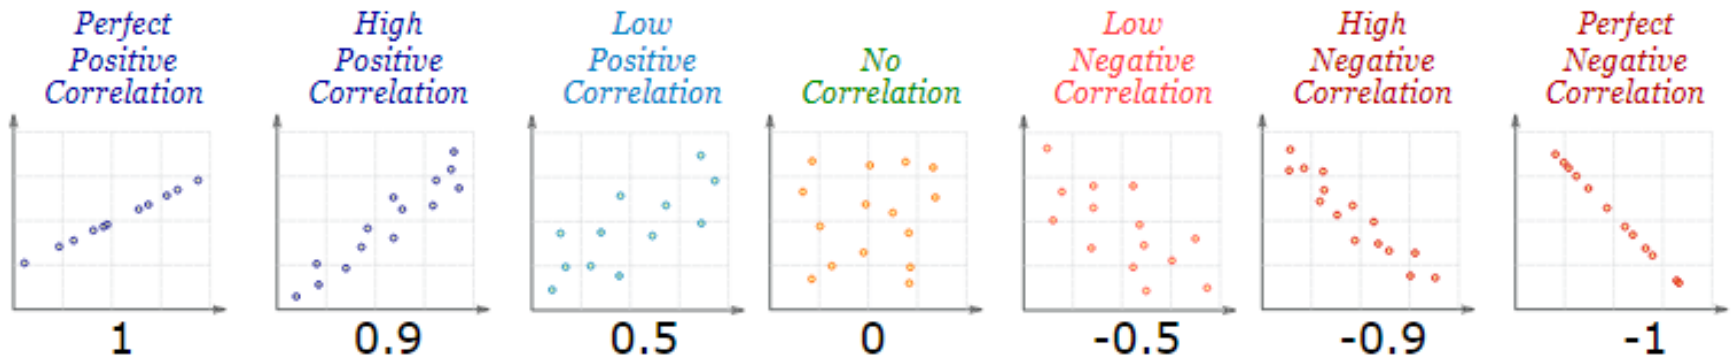
Correlations

- Content

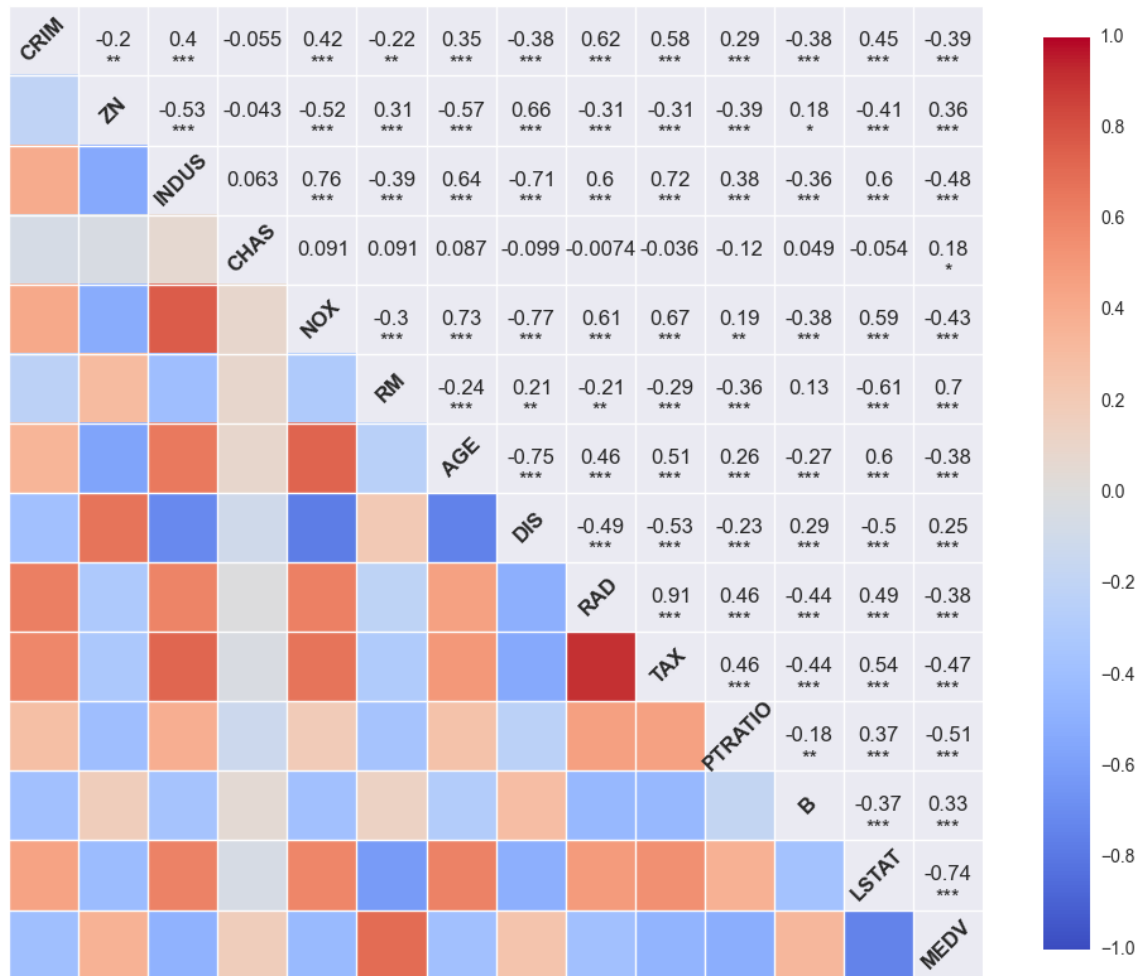


Pearson Correlation

- When two sets of data are strongly linked together, we say they have a High Correlation.
 - Correlation is Positive when the values increase together, and
 - Correlation is Negative when one value decreases as the other increases



Pearson Correlations



Cross Validation

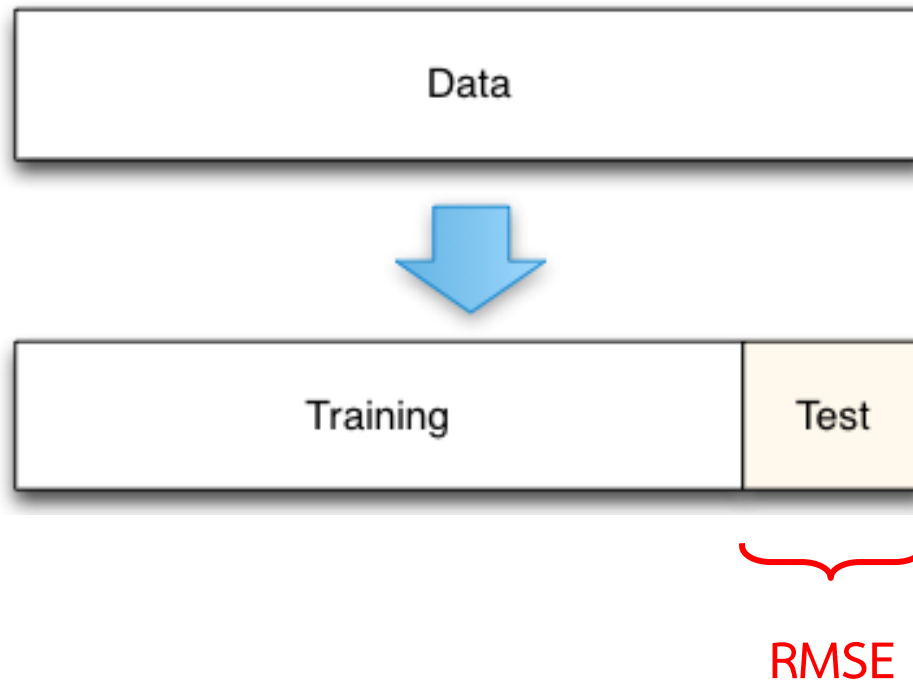


Rotation Estimation

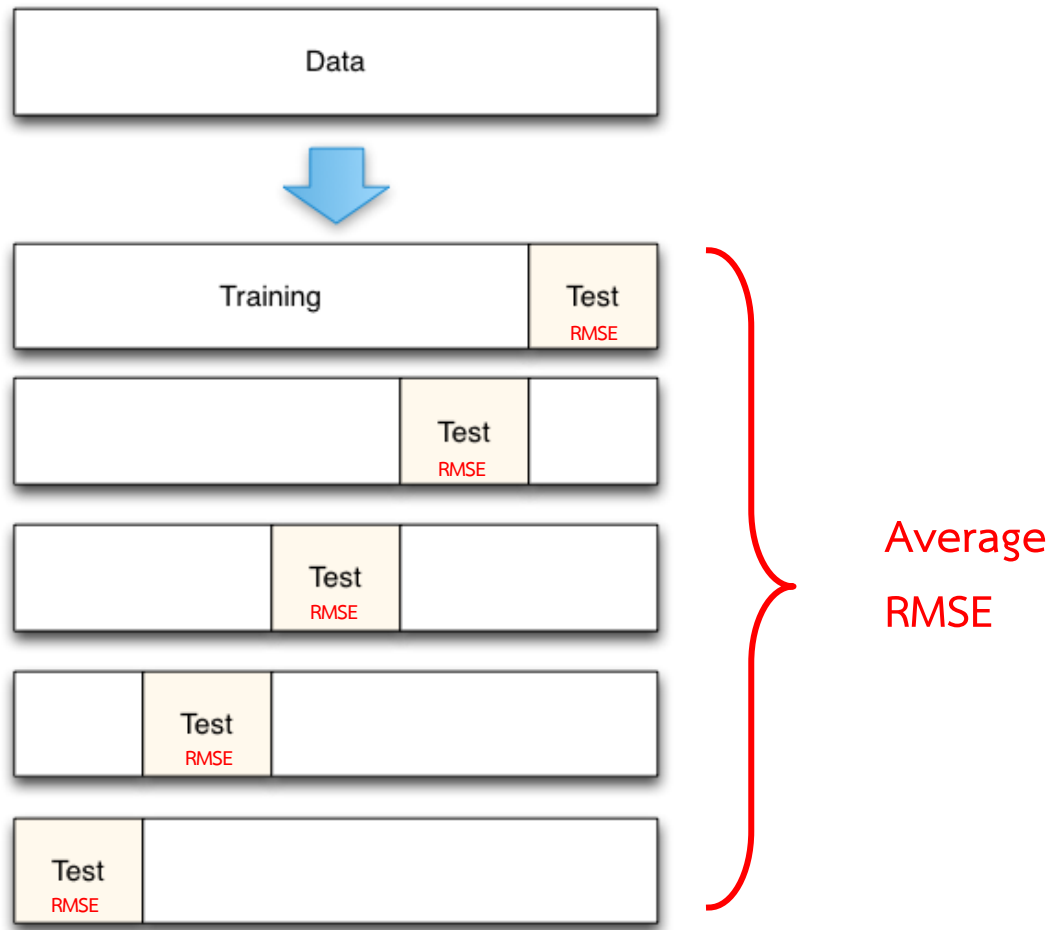
How accurately a predictive model will perform in practice?

- Cross-validation is a technique to evaluate predictive models by partitioning the original sample into a training set to train the model, and a test set to evaluate it.
- In k-fold cross-validation, the original sample is randomly partitioned into k equal size subsamples. Of the k subsamples, a single subsample is retained as the validation data for testing the model, and the remaining k-1 subsamples are used as training data. The cross-validation process is then repeated k times (the folds), with each of the k subsamples used exactly once as the validation data. The k results from the folds can then be averaged (or otherwise combined) to produce a single estimation. The advantage of this method is that all observations are used for both training and validation, and each observation is used for validation exactly once.

Train-Test Split



K-Fold Cross-Validation



Python Code


1. `from sklearn import neighbors, datasets, preprocessing`
2. `from sklearn.model_selection import train_test_split`
3. `iris = datasets.load_iris()`
4. `X, y = iris.data[:, :2], iris.target`
5. `X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=33)`
6. `lm = linear_model.LinearRegression()`
7. `lm.fit(X_train, y_train)`
8. `y_predicted = lm.predict(X_test)`
9. `rmse = mean_squared_error(y_test, y_predicted, squared=False)`

Feature Scaling



Feature Scaling

- Feature scaling is a method used to standardize the range of independent variables or features of data. In data processing, it is also known as data normalization and is generally performed during the data preprocessing step.



Rescaling

- The simplest method is rescaling the range of features to scale the range in $[0, 1]$ or $[-1, 1]$. Selecting the target range depends on the nature of the data. The general formula is given as:

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

- where x is an original value x' is the normalized value. For example, suppose that we have the students' weight data, and the students' weights span . To rescale this data, we first subtract 160 from each student's weight and divide the result by 40 (the difference between the maximum and minimum weights).

Mean Normalization

$$x' = \frac{x - \text{mean}(x)}{\max(x) - \min(x)}$$

- where x is an original value, x' is the normalized value.

Standardization

- In machine learning, we can handle various types of data, e.g. audio signals and pixel values for image data, and this data can include multiple dimensions. Feature standardization makes the values of each feature in the data have zero-mean (when subtracting the mean in the numerator) and unit-variance.

$$x' = \frac{x - \bar{x}}{\sigma}$$

- Where x is the original feature vector, \bar{x} is the mean of that feature vector, and σ is its standard deviation.

Scaling to unit length


- Another option that is widely used in machine-learning is to scale the components of a feature vector such that the complete vector has length one. This usually means dividing each component by the Euclidean length of the vector:

$$x' = \frac{x}{||x||}$$


- In some applications (e.g. Histogram features) it can be more practical to use the L1 norm (i.e. Manhattan Distance, City-Block Length or Taxicab Geometry) of the feature vector. This is especially important if in the following learning steps the Scalar Metric is used as a distance measure.

Box Plot

Before normalization



After normalization



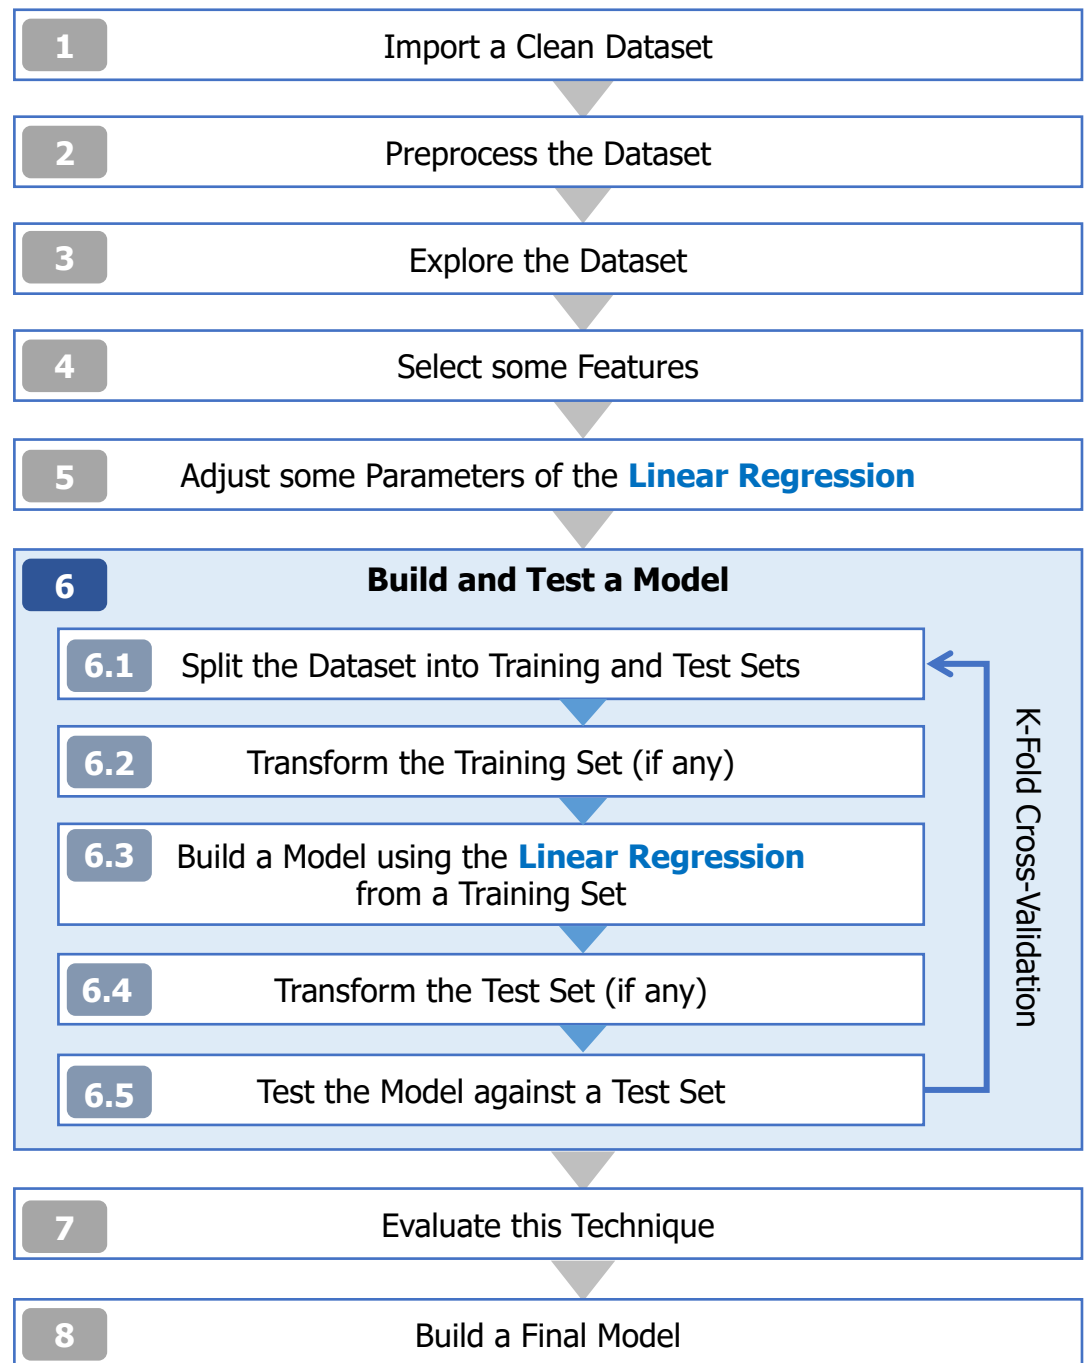
Python Code

1. `from sklearn import neighbors, datasets, preprocessing`
2. `from sklearn.model_selection import train_test_split`
3. `from sklearn.metrics import accuracy_score`
4. `iris = datasets.load_iris()`
5. `X, y = iris.data[:, :2], iris.target`
6. `X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=33)`
7. `scaler = preprocessing.StandardScaler().fit(X_train)`
8. `X_train = scaler.transform(X_train)`
9. `X_test = scaler.transform(X_test)`
10. `# Do Prediction`

Summary



Steps to Create a Linear Regression Model



“

Data is a precious thing and will last
longer than the systems themselves.

”

Tim Berners-Lee